

GAN-Leaks: A Taxonomy of Membership Inference Attacks against Generative Models

Dingfan Chen¹, Ning Yu^{2,3}, Yang Zhang¹, Mario Fritz¹



¹CISPA Helmholtz Center for Information Security, Germany

²Max Planck Institute for Informatics, Germany

³University of Maryland, College Park



Motivation

- Generative adversarial Networks (GANs) have been largely used on privacy sensitive datasets, e.g., face images and medical records
- However, existing works mainly focus on attacks against discriminative models and the privacy risk of generative models have not yet been investigated systematically
- Our work:** Membership Inference Attack against GANs (whether a query sample has been used to train a GAN model?)
- Crucial to understand and control privacy leakage; provides insights for privacy-preserving data sharing

Contributions

- Taxonomy**
 - Categorize attack scenarios against generative models
 - Benchmark future research
- Novel attack models**
 - Generic; easy-to-implement; effective; theoretically grounded
- Extensive evaluation**
 - 3 datasets with diverse data modalities, 5 victim models, 4 attack scenarios ...

Taxonomy

- What information does the attacker know?
 - White-box □/black-box ■?
 - Which GANs' components are accessible? (z : latent code; **Gen**: Generator; **Dis**: Discriminator)
- (1) Full black-box generator
(2) Partial black-box generator
(3) White-box generator
(4) Accessible discriminator (full model)¹
- | | Latent code | Generator | Discriminator |
|--|-------------|-----------|---------------|
| (1) Full black-box generator ^{1,2} | X | ■ | X |
| (2) Partial black-box generator | ✓ | ■ | X |
| (3) White-box generator | ✓ | □ | X |
| (4) Accessible discriminator (full model) ¹ | ✓ | □ | ✓ |

Generic Attack Model

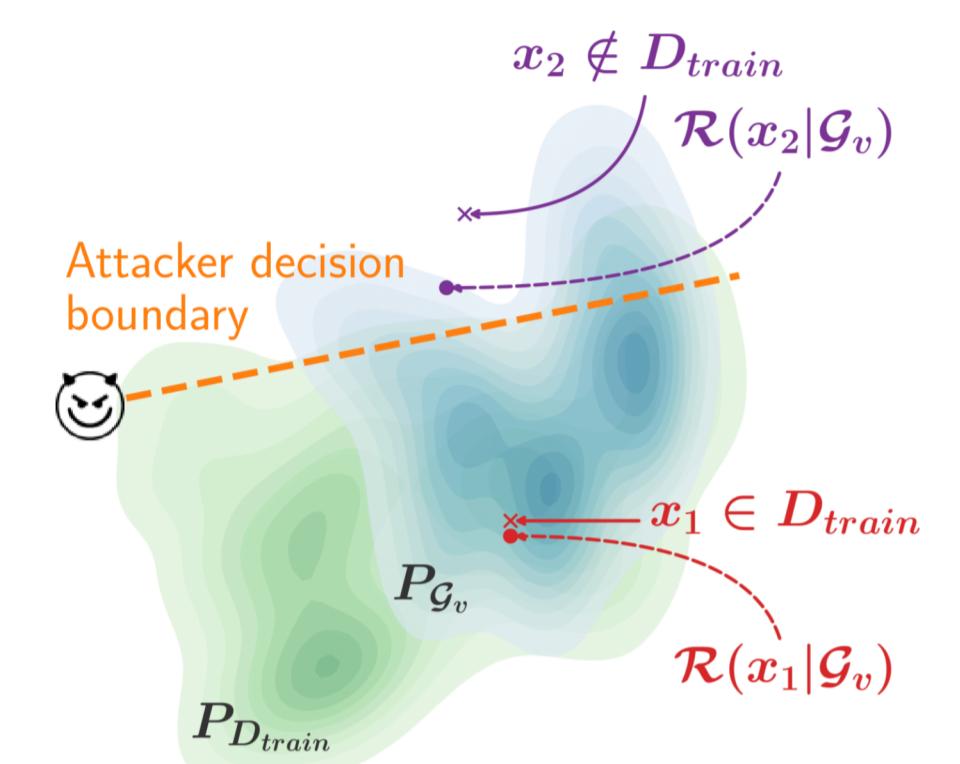
Attacker finds the **best reconstruction** of a query sample given **different types of access** to the victim generator.

- Insight:**
Smaller reconstruction error for training data.

- Generic Model:**
Optimization problem

$$\mathcal{R}(x|\mathcal{G}_v) = \mathcal{G}_v(z^*)$$

$$z^* = \underset{z}{\operatorname{argmin}} L(x, \mathcal{G}_v(z))$$



- Objective:**

$$\underset{z}{\operatorname{minimize}} \quad L(x, \mathcal{G}_v(z)) = \lambda_1 L_2(x, \mathcal{G}_v(z)) + \lambda_2 L_{\text{lips}}(x, \mathcal{G}_v(z)) + \lambda_3 L_{\text{reg}}(z)$$

$$\text{where } L_2(x, \mathcal{G}_v(z)) = \|x - \mathcal{G}_v(z)\|_2^2$$

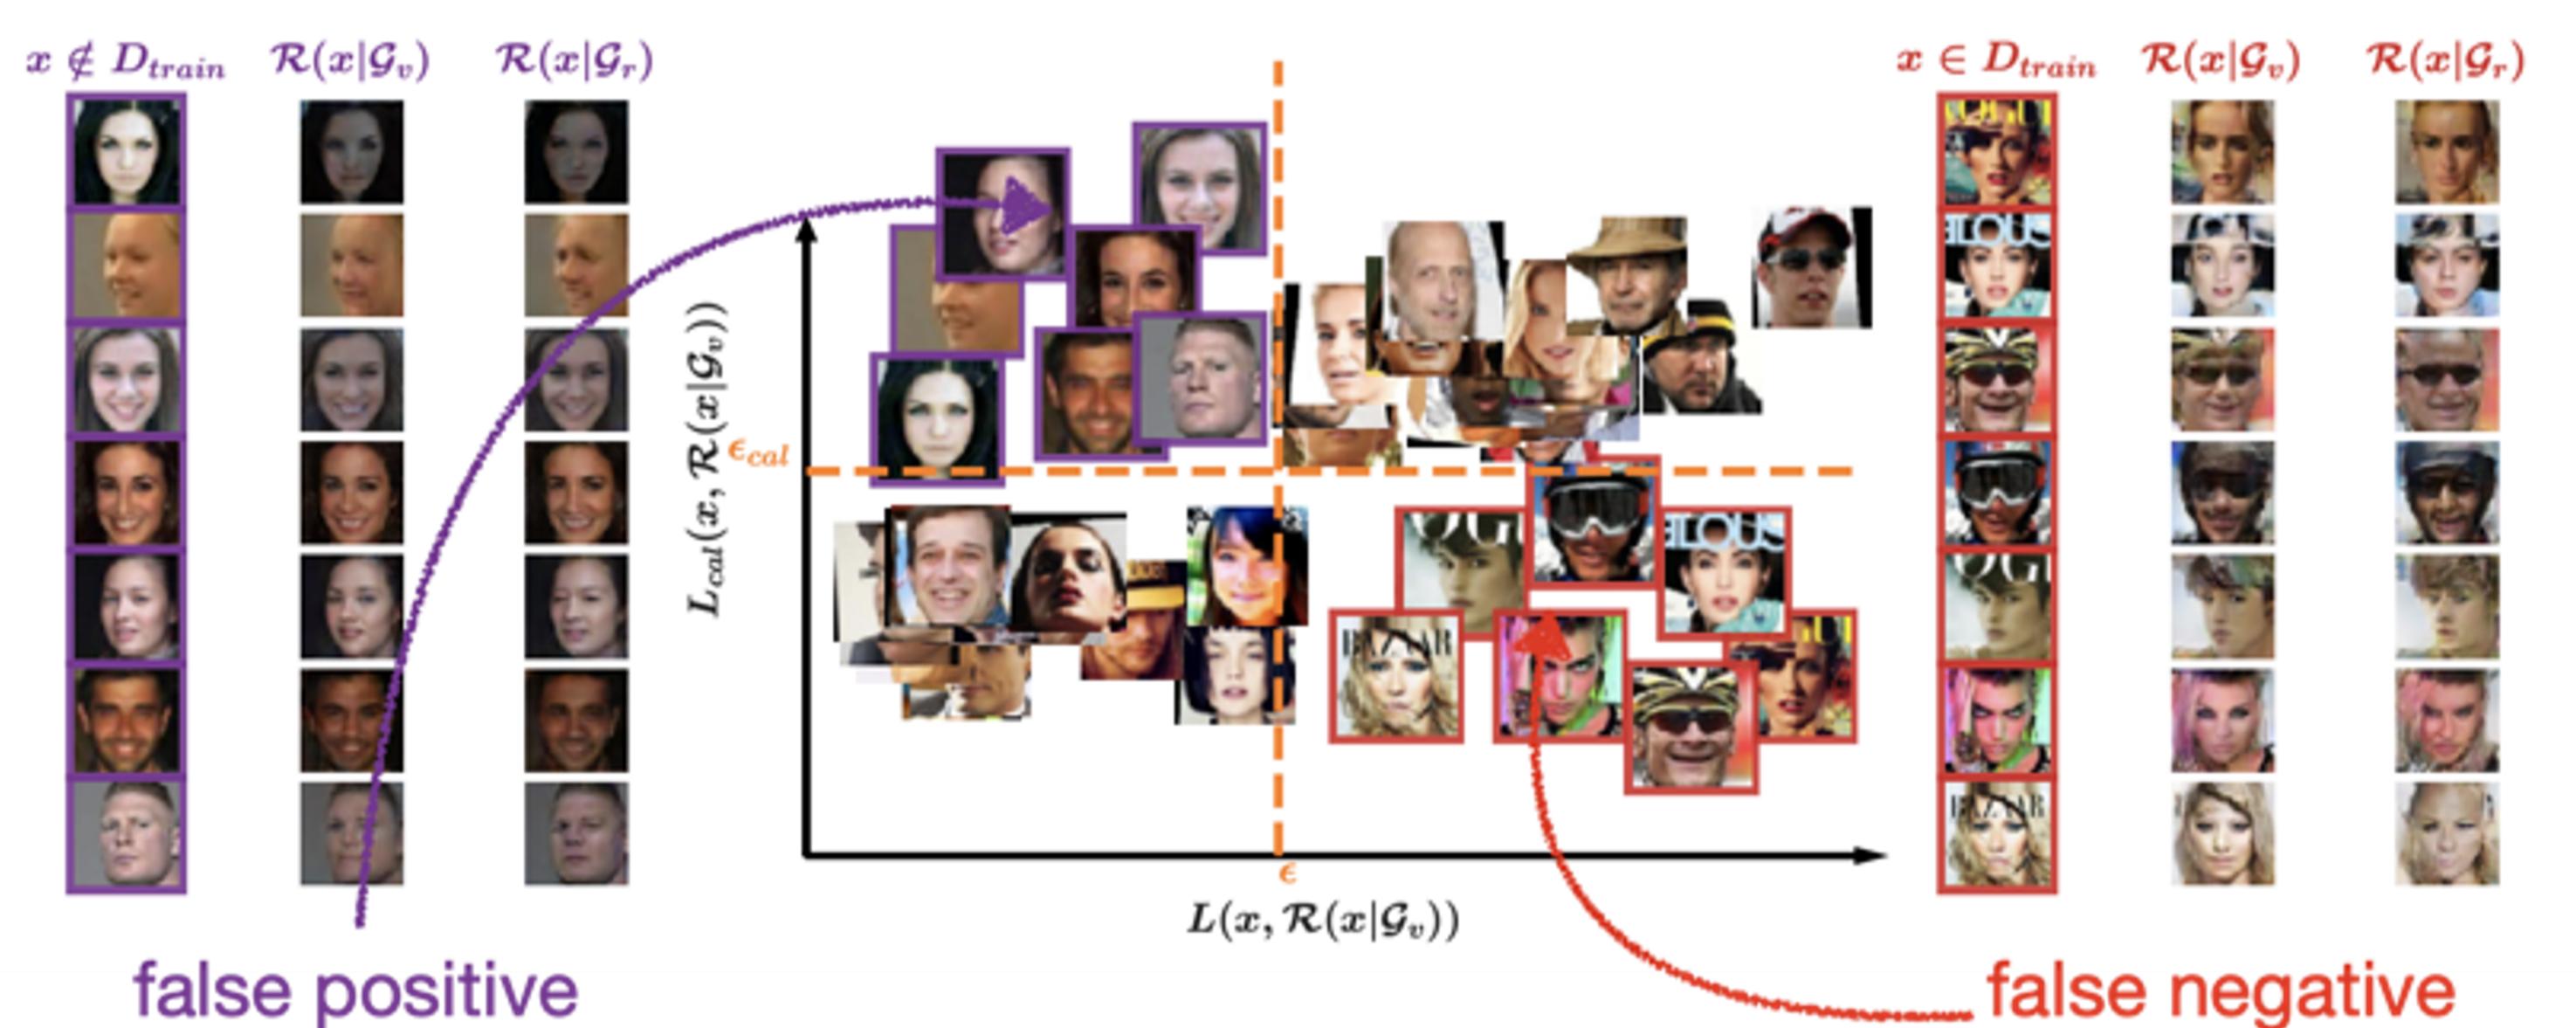
$$L_{\text{reg}}(z) = (\|z\|_2^2 - \dim(z))^2$$

- Different types of access:**
 - (1) Full black-box generator
 - (2) Partial black-box generator
 - (3) White-box generator

KNN search
Powell's conjugate direction method
L-BFGS quasi-Newton method

Attack Calibration

- Problem:**
the reconstruction error is query-dependent ('hard' samples, underrepresented samples)

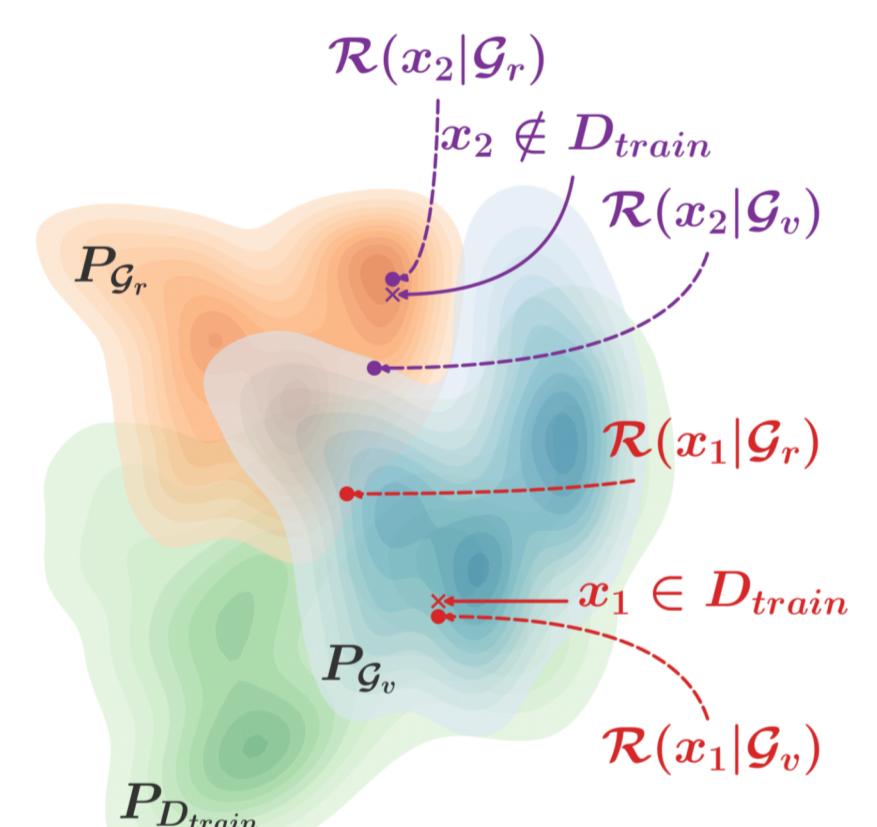


- Solution:** **Attack Calibration**

$$L_{\text{cal}}(x, \mathcal{R}(x|\mathcal{G}_v)) = L(x, \mathcal{R}(x|\mathcal{G}_v)) - L(x, \mathcal{R}(x|\mathcal{G}_r))$$

victim model reference model

- Train a **reference model** with:
 - relevant** but **disjoint** dataset
 - irrelevant** network architecture to victim model
- Theory:** near-optimal under a Bayesian perspective³



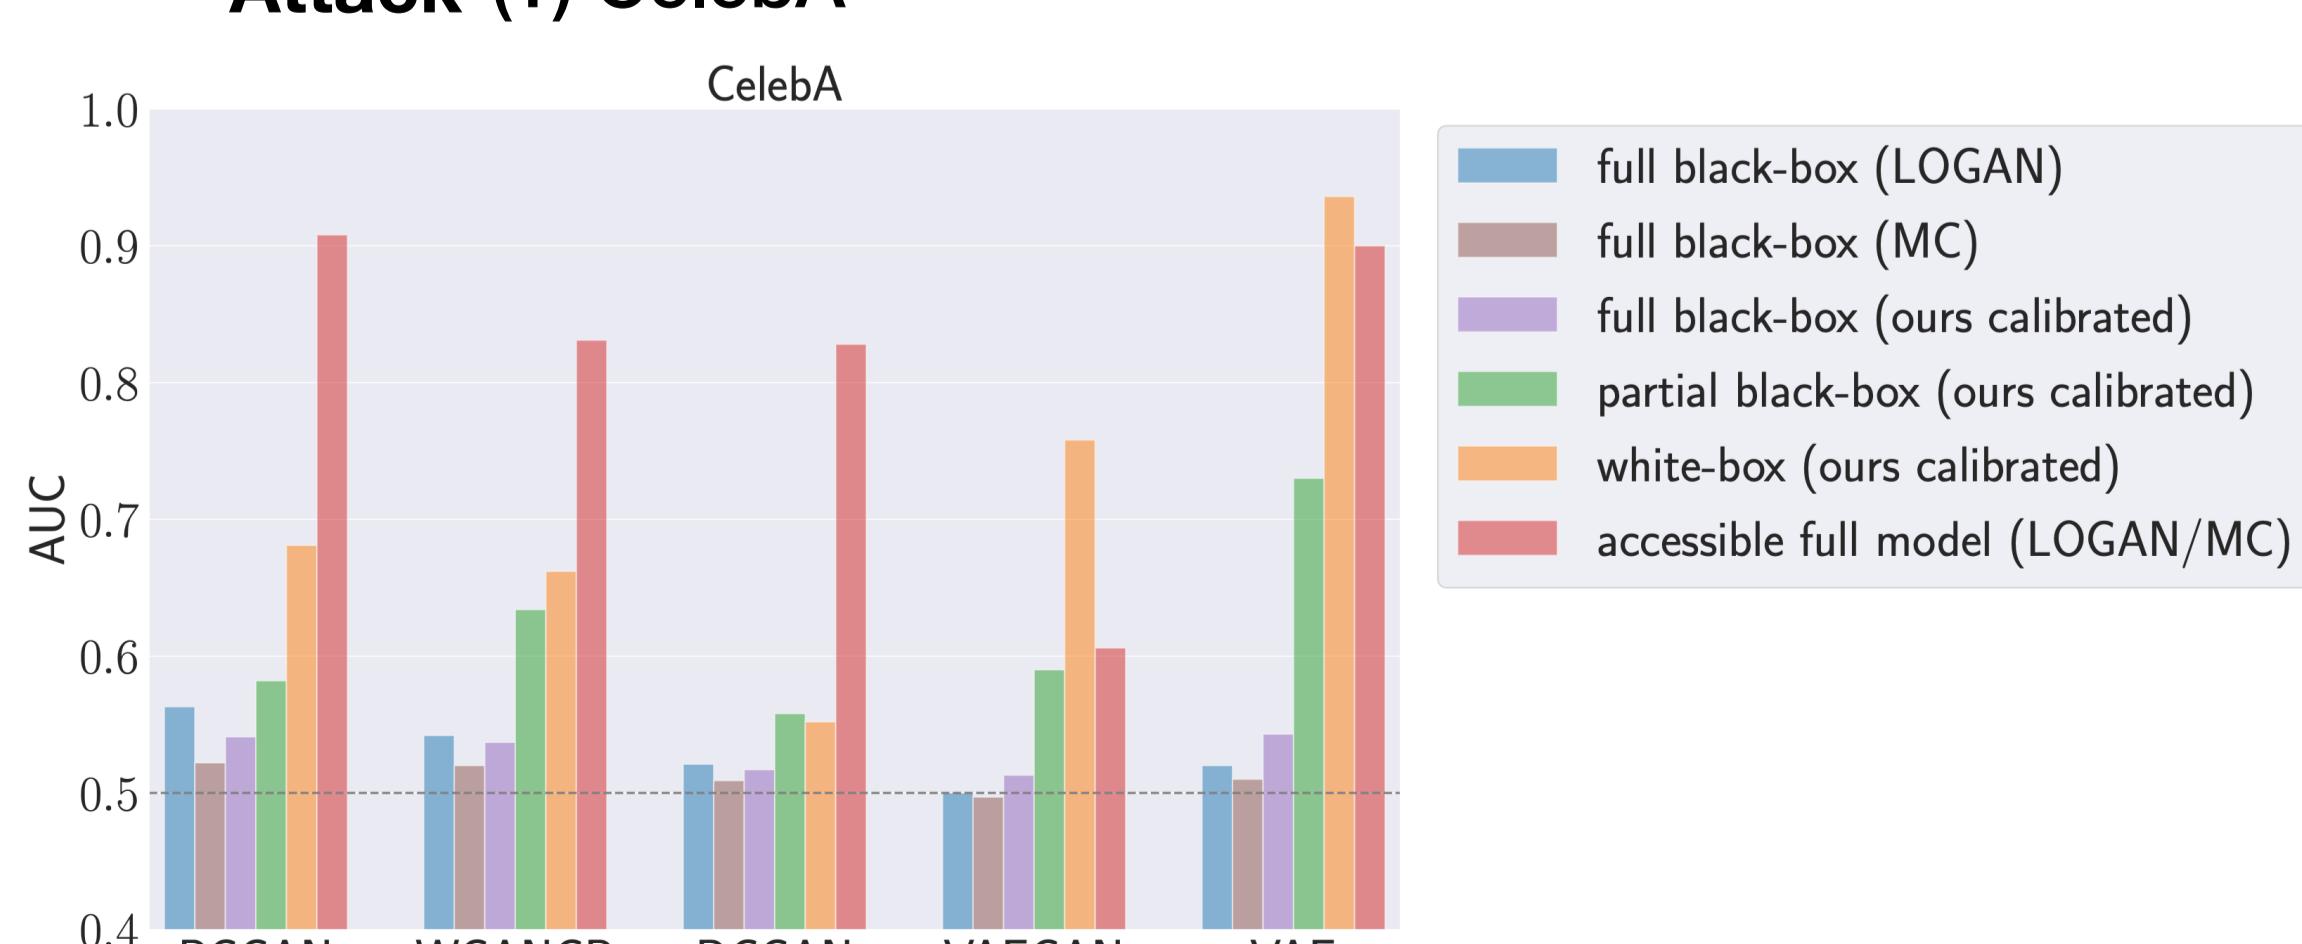
Experiment results

- 3 Datasets:** CelebA (face), MIMIC III (medical), Instagram (location)

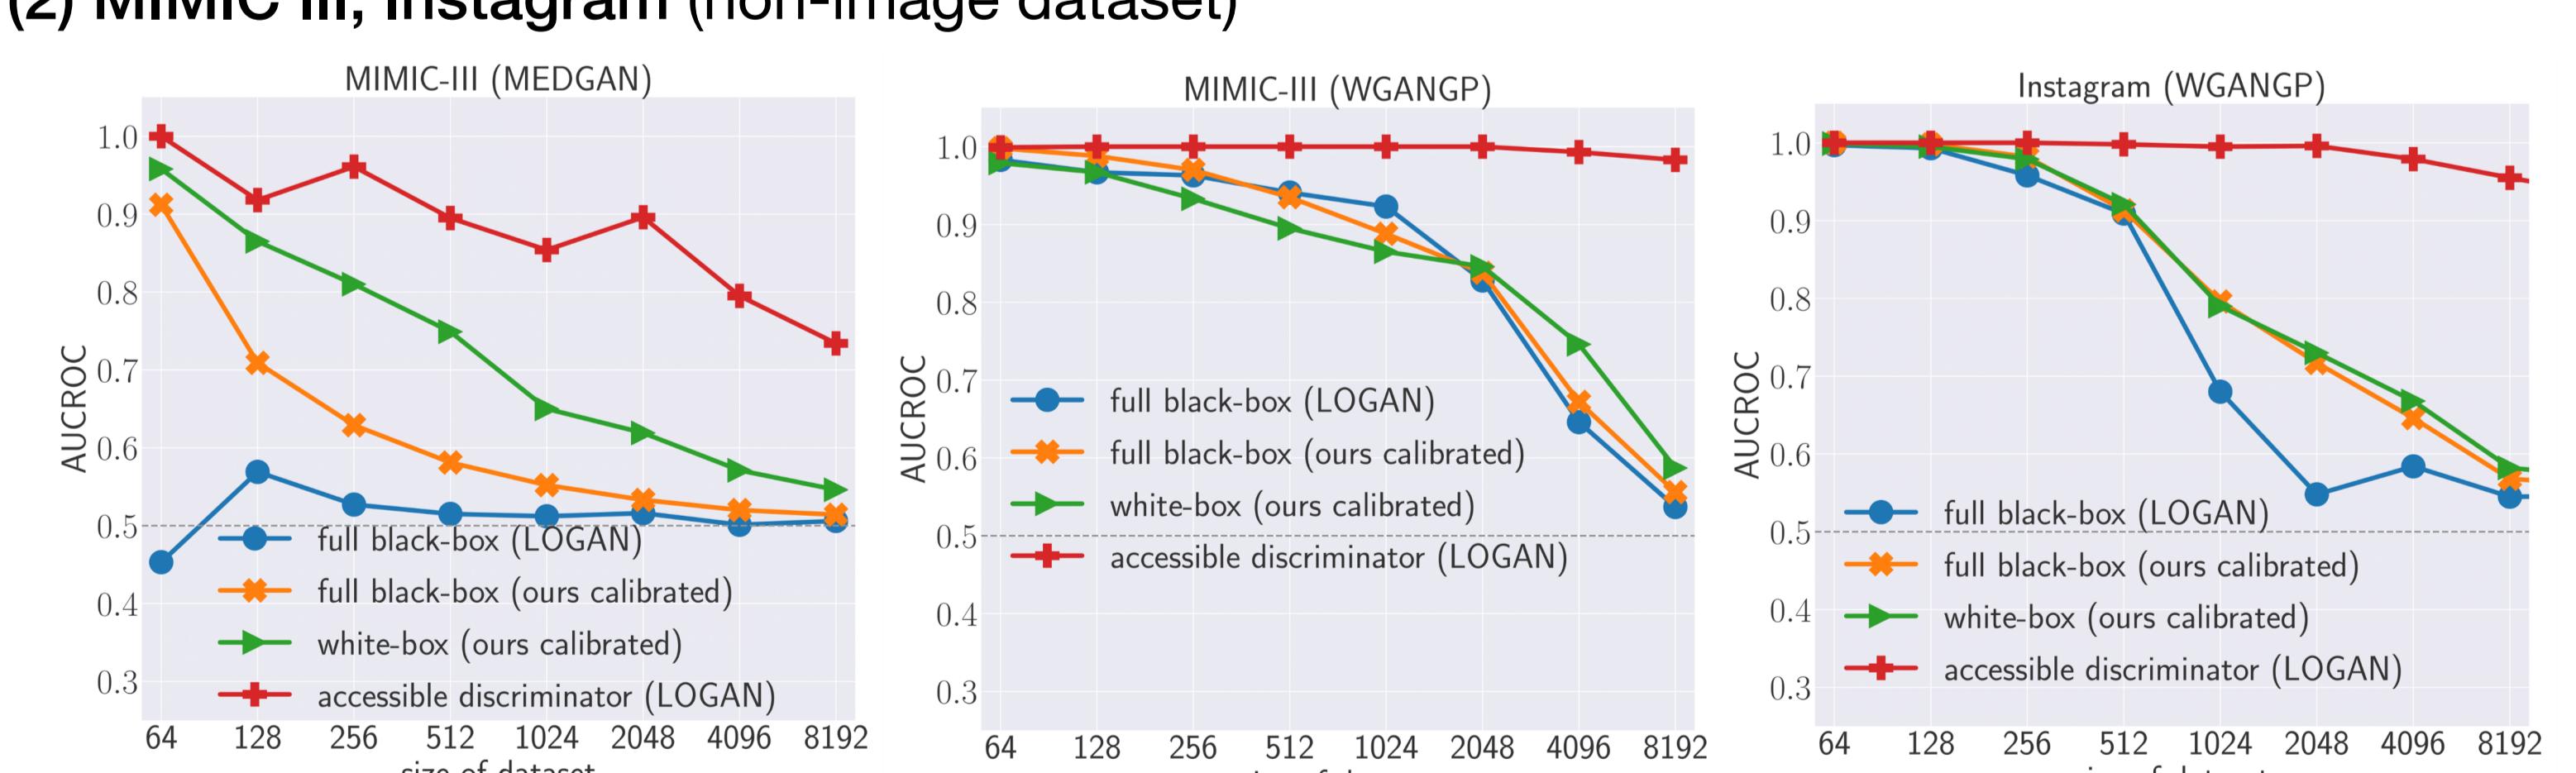
- 5 GAN Models:** PGGAN, WGANGP, DCGAN, VAEGAN, MedGAN

- 2 Baselines:** LOGAN¹, MC²

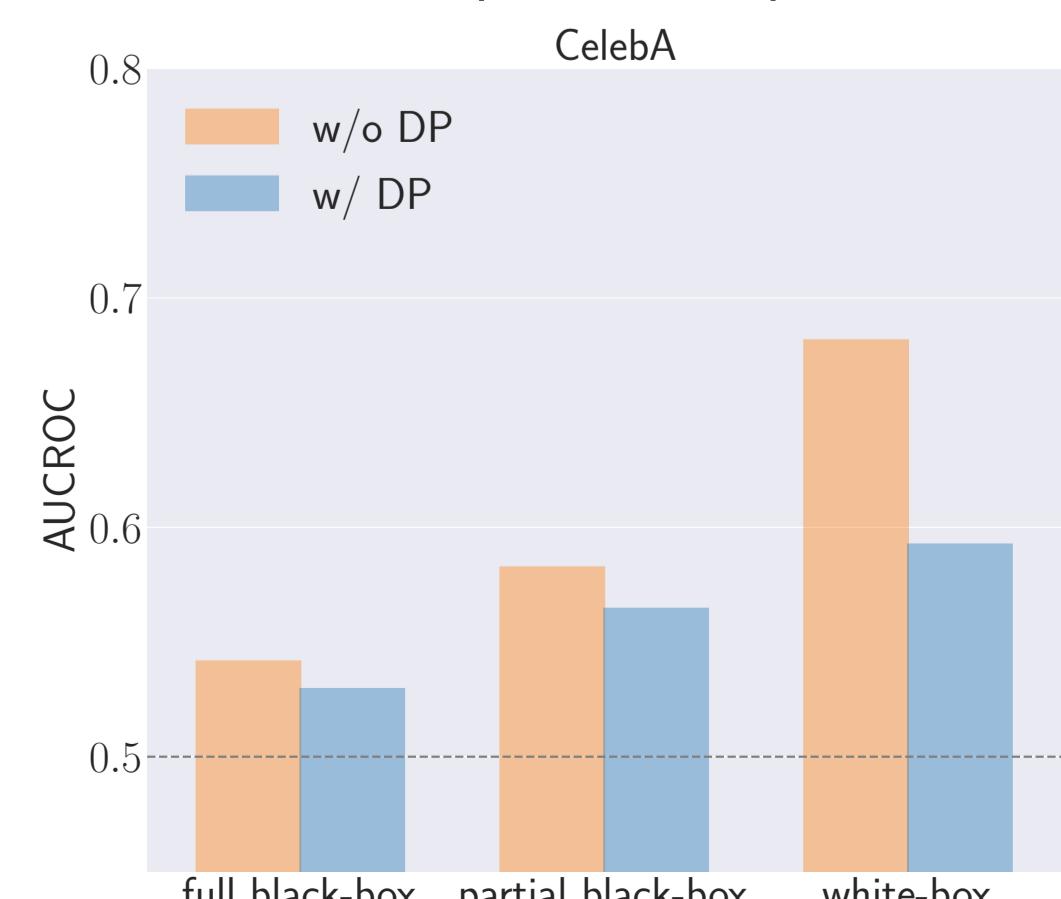
- Results:**
 - Attack (1) CelebA**



(2) MIMIC III, Instagram (non-image dataset)



Defense: (DP-SGD)



Summary

- A simple learning-free attack model works sufficiently well
- Attack performance highly depends on:
 - The size of the dataset
 - Model structure
 - Amount of knowledge about the victim model
- Differential privacy defense is effective against real-world MI attack but compromises utility and efficiency
- Code and models are available on Github: <https://github.com/DingfanChen/GAN-Leaks>

¹ Hayes et al., "LOGAN: Evaluating Privacy Leakage of Generative Models Using Generative Adversarial Networks", PoPETs 2019

² Hilprecht et al., "Monte Carlo and Reconstruction Membership Inference Attacks against Generative Models", PoPETs 2019

³ Sablayrolles et al., "White-box vs Black-box: Bayes Optimal Strategies for Membership Inference", ICML 2019